



# Laziness

Advanced functional programming summer school - Lecture 9

---

Gabriele Keller (& Trevor McDonell, Wouter Swierstra)

# Laziness



# A simple expression

`square :: Integer -> Integer`

`square x = x * x`

`square (1 + 2)`

`= -- magic happens in the computer`

`9`

How do we reach that final value?

# Strict or eager

In most programming languages:

1. Evaluate the arguments completely
2. Evaluate the function call

square (1 + 2)

= *-- evaluate arguments*

square 3

= *-- go into the function body*

3 \* 3

=

9

# Non-strict evaluation

Arguments are replaced as-is in the function body

`square (1 + 2)`

`= -- go into the function body`

`(1 + 2) * (1 + 2)`

`= -- we need the value of (1 + 2) to continue`

`3 * (1 + 2)`

`=`

`3 * 3`

`=`

`9`

# Dopes non-strict evaluation make any sense?

In the case of **square**, non-strict evaluation is worse.

Is this always the case?

`const x y = y` *-- forget about x*

*-- strict*

`const (1 + 2) 5`

`=`

`const 3 5`

`=`

`5`

*- non-strict*

`const (1 + 2) 5`

`=`


`5`

# Sharing expressions

square (1 + 2)  
=  
(1 + 2) \* (1 + 2)

Why redo the work for (1 + 2)?

We can *share* the evaluated result:




# Sharing expressions

square (1 + 2)  
=  
(1 + 2) \* (1 + 2)

Why redo the work for (1 + 2)?

We can *share* the evaluated result:





# Lazy evaluation

Haskell uses a *lazy* evaluation strategy

- Expressions are not evaluated *until needed*
- Duplicate expressions are *shared*

Lazy evaluation never requires more steps than eager evaluation

Each of those not-evaluated expressions is called a *thunk*

# Lazy data-structures

Infinite lists:

```
let xs = 1 : xs  
in take 10 xs
```

```
let ys = [3,6..]  
in take 20 ys
```

```
let primes = ...  
in primes !! 23453
```

Allow to separate generators from filters

# Does it matter?

Is it possible to get different outcomes using different evaluation strategies?

Yes and no

# Does it matter? - Correctness and efficiency

The *Church-Rosser Theorem* states that for *terminating* programs the result of the computation does *not* depend on the evaluation strategy

But

1. Performance might be different

- As **square** and **const** show

2. This applies only if the program terminates

- What about infinite loops?

- What about exceptions?

# Termination

```
loop x = loop x
```

This is a well-typed program

But **loop 3** never terminates

-- *Eager*

```
const (loop 3) 5
```

=

```
const (loop 3) 5
```

=

-- *Lazy*

```
const (loop 3) 5
```

=

```
5
```

Lazy evaluation terminates more often than eager

# Build your own control structure

```
if_ :: Bool -> a -> a -> a
```

```
if_ True t _ = t
```

```
if_ False _ e = e
```

In eager languages, `if_` evaluates both branches

In lazy languages, only the one being selected

For that reason,

In eager languages, `if` has to be *built-in*

In lazy languages, you can build your *own control structures*

# Short-circuiting

```
(&&) :: Bool -> Bool -> Bool
```

```
False && _ = False
```

```
True  && x = x
```

In eager languages, `x && y` evaluates both conditions

- But if the first one fails, why bother?
- C/Java/C# include a built-in *short-circuit* conjunction

In Haskell, `x && y` only evaluates the second argument if the first one is **True**

- `False && (loop True)` terminates

# "Until needed"

How does Haskell know *how much* to evaluate?

By default, everything is kept in a thunk

When we have a case distinction, we evaluate enough to distinguish which branch to follow

```
take 0 _      = []
```

```
take _ []     = []
```

```
take n (x:xs) = x : take (n-1) xs
```

If the number is `0` we do not need the list at all

Otherwise, we need to distinguish `[]` from `x:xs`



# Weak Head Normal Form

An expression is in *weak head normal form* (WHNF) if it is:

- A constructor with (possibly non-evaluated) data inside
  - `True` or `Just (1 + 2)`
- An anonymous function
  - The body might be in any form
  - `\x -> x + 1` or `\x -> if_ True x x`
- A built-in function applied to too few arguments

Every time we need to distinguish the branch to follow the expression is evaluated until its WHNF

# Case study: foldl

```
foldl _ v [] = v
```

```
foldl f v (x:xs) = foldl f (f v x) xs
```

```
foldl (+) 0 [1,2,3]
```

```
= foldl (+) (0 + 1) [2,3]
```

```
= foldl (+) ((0 + 1) + 2) [3]
```

```
= foldl (+) (((0 + 1) + 2) + 3) []
```

```
= ((0 + 1) + 2) + 3
```


# Case study: `foldl`

```
foldl (+) 0 [1,2,3]  
= ((0 + 1) + 2) + 3
```

Each of the additions is kept in a thunk

- Some memory need to be reserved
- They have to be GC'ed after use

# Case study: foldl'



# Case study: `foldl`'

Just performing the addition is faster!

Computers are fast at arithmetic

```
foldl (+) 0 [1,2,3]
= foldl (+) (0 + 1) [2,3]
= foldl (+) 1 [2,3]
= foldl (+) (1 + 2) [3]
= foldl (+) 3 [3]
= foldl (+) (3 + 3) []
= foldl (+) 6 []
= 6
```

# Forcing evaluation

Haskell has a primitive operation to force

`seq :: a -> b -> b`

A call of the form `seq x y`

- First evaluates `x` up to WHNF
- Then it proceeds normally to compute `y`

Usually, `y` depends on `x` somehow

# Case study: `foldl`

We can write a new version of **foldl** which forces the accumulated value before recursion is unfolded

```
foldl' _ v []      = v
foldl' f v (x:xs) = let z = f v x
                    in  z `seq` foldl' f z xs
```

This version solves the problem with addition

# Strict application

Most of the times we use `seq` to force an argument to a function, that is, *strict application*

$$(\$!) :: (a \rightarrow b) \rightarrow a \rightarrow b$$
$$f \$! x = x \text{ `seq` } f x$$

Because of sharing, `x` is evaluated only once



# Profiling

## Something about (in)efficiency

We have seen that Haskell programs:

- can be very short
- and sometimes very inefficient

## Question:

How to find out where time is spent

## Answer:

Use profiling!

# Laziness is a double-edged sword

With laziness, we are sure that things are evaluated only as much as needed to get the result.

But, being lazy means holding lots of thunks in memory:

- Memory consumption can grow quickly.
- Performance is not uniformly distributed.

## Question:

How to find out where memory is spent?

How to find out where to sprinkle **seq**

## Answer:

Use profiling

## Example: `segs`

`segs xs` computes all the consecutive sublists of `xs`.

```
segs [] = [[]]
segs (x:xs) = segs xs ++ map (x:) (inits xs)
> segs [2,3,4]
[[], [4], [3], [3,4], [2], [2, 3], [2,3,4]]
```

This implementation is extremely inefficient.

## Example: `segsinits`

We can compute `inits` and `segs` at the same time.

```
segsinits []          = ([[]], [[]])
segsinits (x:xs)     =
    let (segsxs, initsxs) = segsinits xs
        newinits          = map (x:) initsxs
    in (segsxs ++ newinits, [] : newinits)
segs = fst . segsinits
```